

What Retry Budgets Can Guarantee: Calls, Protected Work, and Partial Repair

ANONYMOUS AUTHOR(S)

A retry limit bounds failed attempts, but its fallback may move every remaining transformation into a critical section. We characterize the call and protected-work guarantees that completion interfaces can provide without knowing the number of concurrent writes. For persistent jobs with mandatory whole-kernel work, a policy using cached completion meets $Q \leq n + r$ in every finite-write execution and minimizes worst protected work simultaneously for every write bound B : the sum of the largest $\min((B - r)_+, n)$ job works. Without that operation, the same universal call cap forces full protection already at $B = r$. Partial repair changes this conclusion. For repeatable write footprints, we characterize exactly when a workflow with a shared retry allowance admits one optimal policy for every budget: each job must have a one-write maximum damage. Per-call costs make precedence and past clean completions relevant; an exact feasibility certificate constructs a common optimum when one exists. These results identify which operations and cost accounting support a claimed retry guarantee. Local and shared retry caps yield different exact frontiers, and charging each call adds a bounded loss to the whole-kernel policy. Java calls realize the abstract bounds under stated restrictions. Roslyn source modifications preserve checked event, snapshot and compiler projections; stronger partial-rebase operations avoid full protected recomputation on eligible updates but retain merge cost. Authored arrival experiments expose conditional background-delay benefits, foreground burst costs and variability, rather than a general latency guarantee.

CCS Concepts: • **Software and its engineering** → **Concurrent programming structures**.

Additional Key Words and Phrases: concurrent programming, retries, resource guarantees, partial repair, online policies

1 Introduction

In optimistic update loops, a caller prepares outside protection and validates before publication. Switching to protected execution after finitely many failures bounds failed calls, but can protect all remaining computation. Which bounds on calls and protected work can be met together, and which completion operations are necessary?

Bounded speculation followed by locking is established practice. With elision enabled, glibc's x86 mutex path defaults to at most three transaction starts before ordinary locking; HotSpot exposes RTMRetryCount for its RTM monitor path [38, 39]. These are motivation for finite attempt budgets, not instances of our theorem: HTM can abort without a write, rolls back partial computation and does not retain an aborted result for our cached completion. We study explicit validate-and-complete interfaces. A caller may cap failures separately for each update, share an allowance across a batch, or require both. These contracts permit different total call counts and protected work. We characterize all three; a global allowance is a specification choice, not a property inferred from those libraries. No latency SLA is assumed.

Consider two dependent transformations with works 8 and 16 and a cap of two validation or completion calls. Every call must complete a transformation. Fresh protected completion performs all 24 units under protection. Cached completion instead prepares outside and, within one completing callback, either reuses the preparation or recomputes after a mismatch. Its worst protected work is 0, 16, or 24 when at most 0, 1, or 2 writes occur (Figure 1). The policy does not receive this bound. The benefit has a cost: at one write its worst total work is 40, compared with 24 for fresh protected completion.

For works (1, 2, 3, 5), a six-call cap allows $r = 2$ failures. Cached completion guarantees at most eight protected work units at four writes. With only fresh fallback, every policy satisfying that




Fig. 1. Operation transitions and the two-job example with $r = 0$. A cached callback completes on either comparison outcome; a cheap mismatch requires a new preparation. Both fixed policies satisfy $Q \leq 2$ in every finite-write environment; neither receives B . Values follow from the model, not timing measurements.

universal cap can be forced to protect all eleven units at two writes. Thus cached completion is necessary for this pair of resource requirements. The units count kernel work, not milliseconds.

The operation contract is consequential. If an update can repair only dirty components, our exact single-job result makes the value of knowing B depend on the allowed write footprints. If each call itself is charged, even a small positive toll can destroy simultaneous optimality. These converses tell a developer when importing the whole-kernel guarantee is invalid. Roslyn’s prepared content-rebase paths supply a concrete example of a stronger operation whose merge cost must be retained.

We make three contributions: exact local/shared call-protection frontiers, a fixed-call-cost loss bound and a suffix feasibility certificate; a footprint-based characterization of when partial repair admits one optimal policy for all budgets, including workflow precedence, per-call costs and policy certificates; and an operational assessment distinguishing Java realizations, Roslyn semantic projections and observed timing. Earlier supplied-price compilation and its full outcomes remain in the artifact. Fallback, online contention control, quantitative synchronization choices and robust dynamic programming are inherited ideas [3, 11, 20, 21, 36, 37]. Our comparison concerns the additional resource equations and their operation-specific lower bounds, not the invention of retry or budget-unaware control.

2 Completion Interfaces and Resource Contracts

Write $\mathbb{N} = \{0, 1, \dots\}$. Jobs $J = \{1, \dots, n\}$, $n \geq 1$, form a finite DAG with predecessor set $\text{Pred}(i)$. Readiness requires all predecessors completed. Completion records an exact immutable output and persists despite later external writes. The objective is one successful foreground transformation per job, not a simultaneous snapshot or transaction. Policies are deterministic and causal; no randomized expected-cost guarantee is claimed.

Each ready job has a pure, opaque, whole kernel receiving its own-input identity and exact completed predecessor outputs. Every execution performs its fixed work $w_i > 0$, whether inside or outside protection. Partial/incremental evaluation, preparation before readiness and discharge by

external work are excluded in this interface. A preparation is tied to its own-input identity and predecessor tuple. Captured outputs cannot be reconstructed by later reads of the mutable map.

Local states are U (unprepared), C (a retained preparation, possibly stale), or D (completed), initially all U; several ready jobs may retain preparations. An invalidating write changes an own-input identity while preserving captured predecessor outputs. Identities never recur. The adaptive environment may place a fresh own-key write immediately before a comparison, including a comparison within a completing callback. It spends at most B successful external writes; unused writes are allowed. The controller does not observe future writes. Budget knowledge, when permitted for a comparator, is stated explicitly.

Let W count all foreground kernel work, L its protected part, and Q validation, publication and completion calls, including failures. Preparations, initialization and external-writer calls are excluded from Q . These are separate resources. Common call overhead is excluded from L . Proposition 4 charges it explicitly in this interface; Section 4.2 treats the distinct repair interface. Let $\Omega_k(T)$ sum the largest $\min(k, |T|)$ works in T . Write $\Omega(T)$ for the total; omitted sets mean J . Set $\Omega_0 = 0$.

Table 1. Whole-kernel operations. Every validation or callback adds one to Q ; preparation adds none. Entries show successor state and added (W, L) . Each callback holds protection only within that operation.

Operation	Start	Match	Mismatch
Prepare	U	C, $(w_i, 0)$	—
Conditional publication	C	D, $(0, 0)$	U, $(0, 0)$
Validation-only callback	C	D, $(0, 0)$	U, $(0, 0)$
Fresh callback	U	D, (w_i, w_i)	—
Cached callback	C	D, $(0, 0)$	D, (w_i, w_i)

Call the first two validating operations *cheap*: they publish on a match and may return without completion on a mismatch. Their label concerns kernel work; it does not assert a zero-time API call. Two modes provide cheap publication and fresh completion; three also provide cached completion. A failed validation leaves the live entry unchanged. Discarding a preparation is free. An inspection is an unprotected read of the current snapshot or identity; it may retain or discard a preparation, but neither publishes nor holds a guard. These reads and local bookkeeping are outside Q . The environment may leave the snapshot unchanged during inspection and write immediately before the next counted comparison, so inspection cannot certify the absence of future interference. Programs can select any ready job after a call, retain preparations and repeat inspections. There is no bare acquire/release, guard-retaining return, cross-job callback or outside completion. All calls and kernels terminate, and foreground scheduling eventually occurs. Noncompletion, including infinite free cycles, has infinite objective. Finite external writing alone does not provide scheduling fairness.

3 Universal Call and Protected-Work Guarantees

For unfinished jobs S , let $\mathcal{A}(S) = \{i \in S : \text{Pred}(i) \cap S = \emptyset\}$ denote the ready jobs.

THEOREM 1 (PROTECTED WORK UNDER A CALL CAP). *In Section 2’s interface, fix integers $B, r \geq 0$. Among programs supplied with B that complete all jobs with $Q \leq n + r$ in every environment with at most B writes, the minimum worst protected work is*

$$L_3^*(B, r) = \Omega_{\min((B-r)_+, n)}, \quad L_2^*(B, r) = \begin{cases} 0, & B \leq r, \\ \Omega, & B > r, \end{cases}$$

with and without cached completion, respectively. Any DAG is permitted.

PROOF. At most r calls may fail to complete. Fresh attempts have disjoint intervals between read and comparison; each failure needs a distinct write. If $B \leq r$ they give $Q \leq n + B$, $L = 0$. Otherwise switch after r failures to cached completion: at most $B - r$ writes remain, giving the stated bound. All-protected gives the two-mode upper bound.

For the two-mode lower bound when $B > r$, reject each cheap call with a fresh own-key write, stopping writes after the $(r + 1)$ st rejection. This uses at most $r + 1 \leq B$ writes. An admissible program cannot incur that many noncompleting calls, so all its jobs complete under protection.

For three modes with $B > r$, target the $k = \min(B - r, n)$ largest jobs. Reject their cheap calls and force their cached calls to recompute by fresh own-key writes immediately before comparison; fresh callbacks already incur w_i inside. Stop writes after $r + 1$ target cheap rejections. Before a next write, at most r such rejections and $k - 1$ cached target completions have occurred, so that write still uses at most $r + k \leq B$ writes. Admissibility excludes $r + 1$ noncompleting calls. Thus every target completes with protected work, giving $L \geq \Omega_k$. Fresh identities invalidate retained preparations, and no inspection or completed callback retains a guard across the chosen boundary. $\square$

THEOREM 2 (ONE POLICY FOR EVERY WRITER BUDGET). *Let $\mathcal{U}_j(r)$ contain the j -mode programs that receive no B and complete with $Q \leq n + r$ in every finite-write environment. Set $U_j(B, r) = \inf_{P \in \mathcal{U}_j(r)} \sup_{\# \text{writes} \leq B} L(P, E)$. Then*

$$U_3(B, r) = L_3^*(B, r), \quad U_2(B, r) = \Omega \mathbf{1}[B \geq r].$$

One policy per mode family attains these values simultaneously for all B and positive works on any DAG, using only readiness and a global cheap-failure counter.

PROOF. Switch after r cheap failures to a fresh read/preparation with cached completion (three) or to fresh protected completion (two). This uses at most r extra calls universally. At $B < r$ no protection is needed; at $B = r$ cached completions after switching still match. Theorem 1 supplies the remaining upper and three-mode lower bounds.

For two modes when $B \geq r$, reject every cheap call with a fresh own-key write until r rejections, then stop. Only protected calls complete before this. Afterwards, another rejectable call could fail in a different finite environment, violating the call cap. Thus all jobs complete under protection with at most $r \leq B$ writes. That additional hypothetical write is absent from the lower-bound execution. Retained preparations and inspections cannot prevent a fresh write before comparison. $\square$

For equal works and $B > r$, the ratio is $n/\min(B - r, n)$, reaching n at $r = B - 1$. The lower bound depends on the universal cap: a budget-aware two-mode policy may use one more cheap call at $B = r$, since that call must match.

3.1 Local Caps and a Shared Cap Can Be Required Together

Suppose job i may incur at most r_i noncompleting calls, and the batch at most $s \leq R = \sum_i r_i$. These caps hold in every finite-write environment. Define the weighted subset value

$$A_c(B) = \max_{T \subseteq J: \sum_{i \in T} c_i \leq B} \sum_{i \in T} w_i,$$

allowing zero-cost items. This is a knapsack value; the contribution is identifying the write cost that activates each job's protection.

PROPOSITION 3 (LOCAL AND SHARED RESOURCE CONTRACTS). *The minimum worst protected work under both caps is*

$$U_2(B; r, s) = \max\{A_r(B), \Omega \mathbf{1}[B \geq s]\},$$

$$U_3(B; r, s) = \max\{A_{r+1}(B), \Omega_{(B-s)_+}\}.$$

One policy attains each curve simultaneously for every B : switch to guaranteed completion when either the current local allowance or the shared allowance is exhausted. Its worst call count is $n + \min(B, s)$.

PROOF. For a local lower bound, target a maximizing subset. Reject each target's cheap calls up to its local cap. In two modes another rejectable call is then universally inadmissible, so the target completes protectively using at most r_i actual writes. In three modes, one additional write invalidates its eventual cached comparison; an earlier fresh completion already pays w_i . Thus each target costs at most $r_i + 1$ writes. The shared lower bounds follow from Theorem 2, since adding local constraints cannot enlarge its policy class.

For the switching policy, if fewer than s failures occur, a protected job has exhausted its own cap: it requires r_i writes in two modes, and $r_i + 1$ in three. If s failures occur, two modes pay at most Ω ; in three, every protected completion still needs its own additional mismatch write, so at most $(B - s)_+$ jobs are protected. These two cases give the maximum above. A fresh preparation after every rejection makes the intervals disjoint. Distributing $\min(B, s)$ cheap failures among the local allowances attains the call bound; its maximum need not coincide with the protection maximum. $\square$

With only local caps take $s = R$, giving $U_2 = A_r(B)$ and $U_3 = A_{r+1}(B)$. Uniform $r_i = r \geq 1$ yields $\Omega_{\lfloor B/r \rfloor}$ and $\Omega_{\lfloor B/(r+1) \rfloor}$, with $Q \leq n + nr$. At $r = 0$ these become Ω and Ω_B , with $Q = n$. Reusing the rejected preparation at the threshold would lose the extra-write requirement of the three-mode bound. A per-job preparation cap is therefore a different contract. Sharing the same total allowance R permits weakly less worst protected work than reserving local allowances; requiring local limits restricts that flexibility.

3.2 What a Fixed Cost per Call Changes

Let every counted call cost $\kappa \geq 0$, and minimize $C = L + \kappa Q$. Let $C^*(B, r)$ be the minimum worst C over the full B -informed program class in Theorem 1. Here $B, r \in \mathbb{N}$ and only the shared cap is imposed; arbitrary local caps are not added to this comparator. Mode choice remains *before* comparison, unlike the repair interface below.

PROPOSITION 4 (A SHARP LOSS BOUND FOR THE THRESHOLD POLICY). *The three-mode global-threshold policy P_r has exact worst cost*

$$C(P_r, B) = n\kappa + \kappa \min(B, r) + \Omega_{(B-r)_+},$$

$$0 \leq C(P_r, B) - C^*(B, r) \leq \kappa \min(B, r).$$

The resulting $r\kappa$ loss bound is sharp for this policy. Exact simultaneous optimality can fail for any positive toll smaller than a job's work.

PROOF. P_r has at most $\min(B, r)$ failures and protection bounded by Theorem 2. Both maxima occur together: reject the first $\min(B, r)$ cheap calls, then, if any writes remain, invalidate cached completions of the heaviest jobs. For every competing program, the resource lower-bound execution has $L \geq \Omega_{(B-r)_+}$ and also $Q \geq n$. Thus $C^* \geq n\kappa + \Omega_{(B-r)_+}$ on that same execution, proving the loss bound.

For one job, $B = r + 1$ lets the environment invalidate every attempt until completion. Immediate completion is optimal at $\kappa + w$, while P_r pays $(r + 1)\kappa + w$. Now let $r \geq 1$ and $0 < \kappa < w$. At $B = 1$ the optimum is 2κ : after its only failure a cheap-first informed policy completes cheaply. On a

write-free prefix through the first counted call, choosing a completing operation can instead be forced to cost $\kappa + w > 2\kappa$. A simultaneous optimum must therefore choose cheap on this prefix. Replay it at $B = r + 1$, reject that call, then invalidate subsequent comparisons until completion. At most r failures and one cached mismatch use at most $r + 1$ writes, forcing $L = w, Q \geq 2$. This costs more than $\kappa + w$. Inspections on other prefixes may reveal past writes; they do not defeat this common write-free prefix. $\square$

The guarantee remains a count of kernel work when fixed overhead is significant. The proposition quantifies its cost sensitivity without identifying a toll with elapsed protected time or claiming minimax regret among all unaware policies.

3.3 Controlling Total Work Along the Protection Frontier

All-cached completion at $r = 0$ uses $Q = n$ and $W = \Omega + L$: avoiding protection after a match requires prior outside work, while a mismatch repeats that work inside. A history can sometimes justify fresh completion of a remaining job, saving its outside preparation while preserving the protection bound. Fix a legal order. The threshold policy freshly prepares before every attempt, uses cheap calls until r failures, then cached completion. The following criterion permits protection-preserving continuations; separate work-order optimization remains in the artifact.

Certifying post-threshold protection choices. After r cheap failures, let h count cached mismatches and ℓ work already protected. Consider a boundary with unfinished jobs S , no retained preparation, and an observed prefix admitting exactly $r + h$ actual writes. This holds for the modeled serial histories; it does not compress histories revealing additional writes.

THEOREM 5 (EXACT SUFFIX FEASIBILITY). *There exists a continuation with at most $|S|$ further calls that meets Theorem 2's protection bound in every finite-write extension iff*

$$C(S, h, \ell) : \quad \ell + \Omega_t(S) \leq \Omega_{h+t}(J) \quad (0 \leq t \leq |S|). \quad (1)$$

PROOF. No further rejectable call is allowed: another adversarial write could force a failure beyond the cap. For necessity, target the t largest remaining jobs. A fresh callback already protects its target; otherwise one fresh own-key write before the eventual comparison invalidates its retained preparations, forcing protected recomputation. At most $u \leq t$ writes force $L \geq \ell + \Omega_t(S)$, while the guarantee requires $L \leq \Omega_{h+u}(J) \leq \Omega_{h+t}(J)$. Inspections cannot retain protection and noncompletion is inadmissible. For sufficiency, freshly prepare/cached-complete in any legal order. Its u mismatches concern distinct jobs, giving $L \leq \ell + \Omega_u(S)$. Every call completes. $\square$

Fresh completion of ready i is therefore feasible exactly when $C(S - i, h, \ell + w_i)$ holds. Cached matching preserves C by subset monotonicity; on mismatch use $w_i + \Omega_t(S - i) \leq \Omega_{t+1}(S)$. These finite inequalities certify allowable actions without B .

The *all-tail* rule finishes all remaining jobs freshly when $\ell + \Omega(S) \leq \Omega_h(J)$; otherwise it continues cached completion. A different rule greedily takes each certified fresh next job.

PROPOSITION 6. *Either rule retains the exact worst- L frontier and universal call cap, and weakly reduces worst W relative to the same-order threshold policy.*

PROOF. At the threshold switch $C(S, 0, 0)$ holds. Certified choices preserve it, every later call completes, and Theorem 2 gives equality. For work, replace each fresh completion in a run by a threshold cached match, retaining every other outcome. This same-order execution has equal W and no additional writes. $\square$

On chain $w = (2, 4, 1, 1)$, $r = 1$, $B = 4$, all-tail reduces worst W from 18 to 17 at $L = 7$, $Q = 5$; both return to 18 at $B \geq 5$. The two certified rules are incomparable (Section 7); feasibility does not establish general W -optimality.

4 When Partial Repair Makes the Budget Valuable

Whole-kernel recomputation makes the first invalidation of a preparation as damaging as further invalidations of that preparation. Some software can instead reuse unaffected components. We now characterize when that operation change destroys a common optimum. We first define a single-job repair interface, then extend it to workflows. This is not a reduction of every Roslyn merge to component work.

Repair interface and observations. One job consists of m components with positive weights w_c . A write changes fresh versions on one nonempty footprint from a fixed finite family Γ . Footprints may overlap and repeat; each occurrence costs one write. Outside protection, preparation reads and computes every component anew. A protected call observes the dirty set D , the union of footprints since that preparation. It either accepts, repairs exactly D and publishes the job, or rejects and prepares again outside. An empty observation can complete with no repair. No partial publication, free dirty-set query, retained protection or cheaper completion operation is available. A fresh whole-job completion costs at least repairing any dirty set. Calls and preparations terminate.

The controller sees D and its history, but neither B nor the number of repeated writes within an interval. It must complete in at most $q \geq 1$ calls for every finite-write execution. The metric in the next theorem is protected repair work; common validation/publication overhead is excluded, not assumed absent. Let

$$G(b) = \max_{0 \leq t \leq b, \gamma_j \in \Gamma} \sum_{c \in \gamma_1 \cup \dots \cup \gamma_t} w_c, \quad W^* = w\left(\bigcup \Gamma\right).$$

Then $G(0) = 0$, G is nondecreasing and subadditive, and $G(m) = W^* > 0$. Unreachable components do not contribute to W^* . Singleton footprints give the prefix sums of weights sorted in decreasing order. A general footprint family need not have a cheap profile: computing these unions is part of the input-processing cost.

THEOREM 7 (EXACT VALUE AND INFORMATION BOUNDARY FOR REPAIR). *For a controller supplied with B , minimum worst repair work is*

$$K_q(B) = G(\lfloor B/q \rfloor).$$

Among controllers receiving no B that attain $K_q(B) = 0$ for every $B < q$, the least possible worst-work envelope is

$$T_q(B) = G((B - q + 1)_+).$$

One policy attains this envelope for all B . For $q \geq 2$, a single controller is optimal for every B if and only if $G(1) = W^$. For $q = 1$, such a controller always exists.*

PROOF. Put $k = \lfloor B/q \rfloor$. The aware policy accepts work at most $G(k)$ and otherwise rejects while a call remains. An observation with greater work requires at least $k + 1$ writes. Rejecting $q - 1$ such observations exhausts all opportunities for a q th: $q(k + 1) > B$. Thus forced acceptance on the last call still meets $G(k)$. This policy also satisfies the call cap outside its supplied budget, where the work guarantee is not asserted. Conversely, a word of at most k footprints realizes $G(k)$. Repeat it after every preparation. Any accepting call pays $G(k)$, and q repetitions use at most $qk \leq B$ writes. Fresh whole-job completion is no cheaper.

For the unaware lower bound, write once in each of the first $q - 1$ preparation intervals. After observation $j < q$, the history is compatible with only j writes in total. Accepting positive damage,

or fresh completion, would violate that budget's zero optimum. Hence the policy must reject all $q - 1$ times. In its final interval spend up to $B - q + 1$ writes realizing $G(B - q + 1)$. No retry remains. The matching policy rejects its first $q - 1$ nonempty observations, then accepts; earlier empty observations complete for zero work.

Every simultaneous optimum meets the preceding zero-work requirements. If $G(1) < W^*$, let $k \geq 2$ be the first index with $G(k) > G(1)$, and set $B = q - 1 + k$. Then $1 \leq \lfloor B/q \rfloor < k$, so $K_q(B) = G(1) < G(k) = T_q(B)$. Conversely, if $G(1) = W^*$, both expressions are zero for $B < q$ and W^* otherwise. For $q = 1$ their arguments coincide. $\square$

This is a condition on the damage a write can cause, not simply the availability of a partial-repair API. If one allowed footprint covers all reachable work, partial repair still admits simultaneous optimality. If writes affect separate components, additional writes can increase the damage and knowledge of B changes the optimal acceptance threshold.

For two unit components, singleton writes and $q = 2$, $K_2(B)$ at $B = 0, 1, 2, 3$ is $0, 0, 1, 1$. Optimality at $B = 1$ forces an unaware policy to reject its first dirty component. Two subsequent writes make both components dirty; its last call costs 2 at $B = 3$. The aware policy at that budget accepts one dirty component immediately and retries only two-component damage. These are complete executions of the specified interface, without assumptions about elapsed time.

Quantifying the loss without a budget. For repair work alone, a finite competitive ratio requires zero work when $K_q(B) = 0$. Theorem 7 therefore identifies the exact optimal ratio as $\sup_{B \geq q} T_q(B)/K_q(B)$. Subadditivity gives an upper bound q : with $k = \lfloor B/q \rfloor$, $B - q + 1 \leq qk$, so $G(B - q + 1) \leq qG(k)$. It is sharp for sufficiently many unit singleton components at $B = 2q - 1$, where the ratio is q . Thus the loss can grow with the call allowance; it is not a constant-price perturbation of whole-kernel optimality.

4.1 A Shared Retry Allowance Across a Workflow

The same boundary extends beyond a single job. Consider n persistent jobs with a DAG, each with its own component/footprint family. A ready job is freshly prepared before each attempted observation; no other preparation is retained. A write affects the currently attempted job, with no cross-job footprint. After a rejection the next ready job may be selected anew. The total call cap is $n + r$, with a shared r ; acceptance pays that job's dirty-component work and permanently completes it. Write $W_i = G_i(\infty)$, $\sigma_i = \min\{b : G_i(b) = W_i\}$, $\Sigma = \sum_i \sigma_i$, and $W^* = \sum_i W_i$. Let $K(B)$ be the known-budget minimum worst repair work for this interface.

THEOREM 8 (WORKFLOW INFORMATION BOUNDARY). *For $r \geq 1$, one policy simultaneously attains $K(B)$ for every B iff $G_i(1) = W_i$ for every job. This holds both for a fixed order and for choices among DAG-ready jobs. Under the condition,*

$$K(B) = \text{sum of the largest } \min((B - r)_+, n) \text{ values } W_i.$$

For $r = 0$, accepting every observation is simultaneously optimal for arbitrary footprints. More generally, the first budget at which $K(B) = W^$ is*

$$B_{\text{sat}} = \Sigma + r \max_i \sigma_i.$$

PROOF. If every job has a one-write maximum, reject the first r nonempty observations globally and otherwise accept. Early acceptances are clean; after r rejections, at most $(B - r)_+$ jobs can accept positive damage. To force the matching lower bound, target that many heaviest jobs, writing once before each target observation to cause maximum damage. A target completion costs W_i , and at most r rejections add writes, so $r + k \leq B$ writes suffice for k targets. An attempted $(r + 1)$ st

rejection would itself violate the call cap within budget; truncate the adversary there. Selection and precedence do not avoid completing the targets.

For necessity, all $K(B)$ with $B \leq r$ are zero. A simultaneous optimum must therefore reject when one write is issued before each of its first r observations: after observation t , the history is compatible with only $t \leq r$ writes in total. No job has yet completed. Now spend σ_i writes to saturate each job on its next attempt. No rejection remains, so the policy pays W^* with $B_0 = \Sigma + r$ writes. If some $\sigma_j \geq 2$, a competing policy reserves all r rejections for fully dirty observations of job j and accepts everything else. Full total work would require $\Sigma + r\sigma_j > B_0$ writes, as every invocation uses a fresh preparation. Its worst cost is therefore strictly below W^* ; the attainable cost set is finite. This contradiction proves necessity even in a prescribed order.

At $r = 0$, every call must complete and acceptance dominates fresh whole-job completion. The value is $\max_{\sum b_i \leq B} \sum_i G_i(b_i)$ independently of order. For the final statement, saturating every attempted observation uses at most $\Sigma + r \max_i \sigma_i$ writes on an admissible run and forces full work. Below that budget, reserve all r rejections for full observations of a maximum- σ_i job. Full work would require exactly that many writes, so the worst cost is strictly smaller. $\square$

For two jobs, one with two unit singleton components and one with a single unit component, $r = 1$ gives $K(B) = 0, 0, 1, 2, 2, 3$ at $B = 0, \dots, 5$ in either order. Yet optimality at $B = 1$ forces a first rejection; three subsequent writes saturate both jobs and force cost 3 at $B = 4$, where $K(4) = 2$. The shared allowance does not remove the information gap. The theorem identifies its existence and saturation threshold; it does not give a closed form for every intermediate value in an arbitrary nonsaturating workflow.

4.2 Charging Every Call Changes the Boundary Again

Repair work omits the protected cost of validation itself. To expose this assumption, charge $\kappa \geq 0$ on *each* call and a common $\mu \geq 0$ once upon completion. Acceptance on call t costs $\mu + t\kappa + w(D)$. Preparation remains outside this metric. Let M be the sorted multiset of qm numbers

$$\{t\kappa + G(k) : 1 \leq t \leq q, 0 \leq k < m\},$$

with one-based entries M_j and repeated entries retained.

THEOREM 9 (TOLL CURVE AND SIMULTANEOUS-OPTIMALITY BOUNDARY). *The known-budget optimum is $\mu + C_q(B)$, where*

$$C_q(B) = \begin{cases} \min\{\kappa + W^*, M_{B+1}\}, & B < qm, \\ \kappa + W^*, & B \geq qm. \end{cases}$$

For $q \geq 2$, simultaneous optimality holds exactly in the following cases: $\kappa = 0$ and $G(1) = W^$, or $\kappa \geq W^*$. It fails for every $0 < \kappa < W^*$. The common cost μ does not change these conditions. At $q = 1$, first acceptance is always simultaneously optimal.*

PROOF. Remove μ and fix $T < \kappa + W^*$. On call t , an unacceptable dirty set has $t\kappa + w(D) > T$. Its least write cost is $a_t(T) = \min\{k : t\kappa + G(k) > T\}$. Such an index exists since $G(m) = W^*$. A sequence of q unacceptable observations is possible exactly when $\sum_t a_t(T) \leq B$: each interval can independently realize a minimum-cost witness. Otherwise accepting the first acceptable observation finishes within q calls and cost T . Because G is nondecreasing, $a_t(T)$ counts that shifted list's entries at most T . Hence $\sum_t a_t(T)$ is the rank of T in M , with ties. The smallest achievable threshold below the first-accept cap is M_{B+1} , proving the curve.

If $\kappa > 0$ and a policy ever rejects a first observation, follow it by maximum damage until it accepts. At a sufficiently large finite budget it pays at least $2\kappa + W^*$, exceeding the aware first-accept

cap $\kappa + W^*$. Thus a simultaneous optimum must accept every first observation. Let k^* be the first index with $G(k^*) = W^*$. With known budget k^* , accept all first observations except maximum damage; reject that one and the next observation is necessarily clean. Its worst cost is at most $\max\{\kappa + W_<, 2\kappa\}$, where $W_< < W^*$ is the largest attainable nonmaximum damage. For $0 < \kappa < W^*$ this is strictly smaller than the first-accept value $\kappa + W^*$. Conversely, when $\kappa \geq W^*$, presenting a first observation with damage $G(B)$ forces either $\kappa + G(B)$ on acceptance or at least 2κ on rejection. First acceptance attains the former, which is no larger. The zero-toll case is Theorem 7. $\square$

The zero-toll guarantee is therefore sensitive to the objective: even an arbitrarily small positive per-call charge precludes a common optimum until that charge reaches the maximum repair work. This does not refute a bound that deliberately measures kernel work alone. It identifies which overhead can be ignored for that particular optimality claim.

The supplement also compiles the best deterministic competitive acceptance threshold. If $c(D)$ is the least footprint count realizing D and e sums those counts after rejections, acceptance on call t is safe for ratio R iff $\mu + t\kappa + w(D) \leq R(\mu + C_q(e + c(D)))$. The comparator is smallest at this least-consistent budget. Monotonicity reduces ratio feasibility to a minimum-write failing path, with exact rational upper/lower certificates. Given correct profiles, the specialized compiler uses $O(qm \log(q+1) + qm \log U)$ arithmetic operations and $O(qm + q)$ space, where $U = \mu + q\kappa + W^*$ after integer scaling. Arbitrary footprint profiling and runtime minimum-cover computation remain separate costs; singleton footprints use sorting and $c(D) = |D|$.

4.3 With Call Charges, Precedence and History Matter

Return to the workflow interface of Section 4.1, charging $\kappa > 0$ per call. Write $K^\kappa(B)$ for its known-budget value, $H(B) = \max_{\sum b_i \leq B} \sum_i G_i(b_i)$ and $C_{\max} = n\kappa + \sum_i W_i$.

THEOREM 10 (CALL CHARGES IN WORKFLOWS). *For independent jobs and $r \geq 1$, a simultaneous optimum exists iff $\kappa \geq \max_i W_i$. This condition is sufficient on every DAG: first acceptance has exact value $n\kappa + H(B)$. On a general DAG, a sink with $W_i > \kappa$ prevents simultaneous optimality. The condition on all jobs is no longer necessary.*

PROOF. For sufficiency, fix an allocation attaining $H(B)$. Present its dirty set on each job's first attempt and leave that job clean after a rejection. Accepting pays the allocated damage; rejecting incurs an extra toll at least as large. This matches the first-accept upper bound, regardless of order.

For necessity, a simultaneous optimum cannot reject after a prefix whose completed jobs were all fully dirty: saturating every remaining attempt would then force cost above $C_{\max}$ at some finite budget. Hence fully dirty first observations force it to pay $C_{\max}$ with Σ writes. If a job j with $W_j > \kappa$ can be placed last, an aware policy at budget Σ accepts preceding jobs. If any observation is less than full it accepts throughout, costing strictly below $C_{\max}$. Otherwise all Σ necessary writes are spent; rejecting j once then completing cleanly costs $C_{\max} - W_j + \kappa < C_{\max}$. The finite outcome set makes the worst cost strictly smaller. Every independent job, and every DAG sink, can be last.

For the final claim, let A have singleton components of weights 1 and 2, let C have one component of weight 1, require A before C , and take $r = 1, \kappa = 2$. First acceptance has curve 4, 6, 7, 8 at $B = 0, 1, 2, \geq 3$. Matching lower bounds subtract the mandatory toll 4. At $B = 1$, dirty A 's weight-2 component: acceptance or rejection costs at least 2. At $B = 2$, do the same, then dirty C after acceptance (cost 3), or the weight-2 component again after rejection (cost 4). At $B = 3$, first saturate A with two writes; then dirty C after acceptance (cost 4), or the weight-2 component after rejection (cost 4). Thus first acceptance is a common optimum although $\kappa < W_A$. Reversing the order makes A a sink and removes every common optimum. $\square$

History-dependent rejection can be essential even with one-component jobs. Consider chain (a, h, a) , $h > a > 0$, with $r = 1$, $\kappa = a$. Accept throughout except that a dirty middle job is rejected once if the first job completed cleanly. Its exact curve is $3a, 4a, 4a + h, 5a + h$ at $B = 0, 1, 2, \geq 3$. First acceptance instead costs $3a + h > 4a$ at one write. For upper bounds, a retry leaves at most $B - 1$ writes; whenever its extra toll a is paid, the clean first job has saved a . For matching lower bounds, one dirty observation forces at least a beyond baseline. With two writes, leave the first job clean and dirty the middle, then dirty the last after acceptance or the middle again after rejection, forcing $a + h$. With three writes, dirty the first job: after acceptance apply that two-write strategy to the last two jobs; after rejection dirty the middle and last instead. This forces $h + 2a$. A preceding clean rejection already spends the toll and leaves no retry. No writer budget is supplied; $(a, h) = (1, 3)$ gives the saved seven-case separation.

Constructing a Common Policy for Other DAGs. A finite recurrence decides whether a common optimum exists beyond these closed-form cases. Set $E = \Sigma + r \max_i \sigma_i$; saturating every attempt proves $K^\kappa(b) = C_{\max}$ for $b \geq E$. Let e be the least-consistent write count clipped to E , t the remaining rejections, and a the cost already paid. Define $\Theta(\emptyset, t, e) = K^\kappa(e)$ and, with $e_D = \min(e + c(D), E)$,

$$\Theta(S, t, e) = \max_{i \in \mathcal{A}(S)} \min_D \max \{ \Theta(S - i, t, e_D) - \kappa - w_i(D), \Theta(S, t - 1, e_D) - \kappa \}, \quad (2)$$

omitting the rejection term at $t = 0$. A continuation exists exactly when $a \leq \Theta(S, t, e)$; in particular, $\Theta(J, r, 0) \geq 0$ decides common optimality.

To prove this, consider a completed history and sum the minimum footprint-word lengths of its observations. Fresh preparations and independent own-job writes realize those words in separate intervals, so the sum is the least budget consistent with that history. Every larger budget permits the same history by leaving writes unused. Since K^κ is nondecreasing, meeting its value at the least budget is necessary and sufficient for that history to meet every consistent target. Above E all targets equal $C_{\max}$; clipping this coordinate is exact even though a partial observation's minimum word can exceed its job's full-saturation count. Accumulated cost a is never clipped.

Induct on $|S| + t$. With no jobs left, feasibility is the interval $a \leq K^\kappa(e)$. For a fixed ready job and observation, accepting adds $\kappa + w_i(D)$ and rejecting adds κ ; either strictly decreases the induction measure. Each child's feasible costs therefore give a shifted downward interval. The choice after observing D takes the larger endpoint. Requiring success for every observation takes the minimum; choosing a job before that observation takes the maximum. This proves Equation (2) with its required $\exists i \forall D \exists \text{action}$ order. A maximizing job and any satisfied child inequality supply an executable policy using history alone. If the initial interval excludes zero, the failed alternatives identify an adversarial finite observation path for each attempted controller choice.

For each minimum footprint count, only its maximum attainable damage is needed in the adversarial alternatives. With I unfinished sets and $h = \max_i (m_i + 1)$, the table has $O(I(r+1)(E+1))$ entries and uses $O(In(r+1)(E+1)h)$ arithmetic operations, given profiles and the comparator. For a fixed order, $I = n + 1$ and there is one job choice. This removes accumulated-cost enumeration, not potentially exponential DAG states or footprint profiling. It is a specialization of finite minimax and information-state reasoning, not a new general dynamic-programming principle.

5 Actual External Writes in Java

Use Java 17's `ConcurrentHashMap` with one foreground thread, distinct persistent keys, immutable identity-equal wrappers and fresh, never-reused identities for external writes. Each job updates only its fixed key. Keys are not removed; results are non-null; kernels use own snapshots and captured parents and terminate without exceptions. Conditional replacement accepts identity or

equals equality [26, 27]; atomic compute invokes its callback once, under a bin monitor in the inspected source [25, 27]. Different keys may share bins.

Implement Table 1 using `replace(key, expected, result)` and `compute`. Validation-only callbacks return the prepared result on match, otherwise the unchanged entry with a failure flag; cached callbacks recompute on mismatch, and fresh callbacks always compute under protection. Capture permanent milestones locally rather than rereading the map.

Callbacks perform no other-key map operations, other foreground kernels or direct guard work, including on cached matches; no guard survives. Multiple preparations are allowed. Bin collisions add no unperformed kernel work. Both `replace` and `compute` can acquire a bin monitor: L counts kernel work within it, while Proposition 4 separately accounts for fixed call costs. Assume scheduling and call completion: finite write counts do not imply fairness.

THEOREM 11 (RESOURCE BOUNDS FOR ACTUAL WRITES). *Under the stated API, monitor and scheduling assumptions, Theorems 1–2’s exact call/protected-kernel-work bounds hold for Java programs using the listed operations, with separate or shared internal bins. This is a mathematical refinement argument for those operations, not a mechanically verified property of an entire JDK binary.*

PROOF. For the upper bounds, the serial policies freshly read and prepare after every cheap failure. Their read-to-comparison intervals, including cached completions, are disjoint; every mismatch contains a distinct successful write. Thus at most B mismatches occur. Every fresh/cached completing call completes exactly one job, and a cached mismatch performs precisely that job’s kernel inside.

For the lower bounds, use the target-job adversaries of Section 3. Immediately before a chosen comparison acquires its monitor, a successful own-key write installs a fresh identity, invalidating every retained preparation for that job. A cheap call fails; a cached call recomputes inside; a fresh callback already incurs its whole kernel inside. Stale copies cannot complete cheaply, captured parent outputs persist, and no monitor survives the preceding call. Bin collisions therefore do not prevent the next chosen write. The prohibition on other-job work within a callback ensures it cannot complete several targets for one call. The same write counts and hypothetical extra-write argument apply. $\square$

Operational accounting. Native kernels fold captured predecessor seeds and an immutable array of length $8w_i - |\text{Pred}(i)| > 0$, executing exactly $8w_i$ selected operations. Counters record their protected subset and the specified calls separately. Supplementary scalar tests add $8p_i$ per protected kernel, $8v_i$ per conditional call and $8k_i$ per callback. Allocation, monitor waiting, initialization, background work and other API instructions are excluded; charges are not inferred from timing.

6 Source Integrations and Their Boundaries

Roslyn document updates without a writer budget. The `Workspace.cs` loop at commit 6c4a46a3 transforms a captured immutable solution outside a semaphore, validates identity under it and retries mismatches. Publication preserves notification ordering [4, 5]. We modify only the `SourceText` path. After r global failures, two-mode completion transforms the latest solution inside; three-mode completion prepares outside, then reuses on match or recomputes inside on mismatch. The selector receives no writer/price tables.

Replayability and safety under protection are separate obligations. The scratch document list must be cleared each attempt; source review discusses replay and event ordering [5]. Our initial two-mode patch deadlocked on same-text notification reentry; restoring the pre-lock identity no-op repairs it. Four later missing-document cases timed out when virtual `GetDocumentName` reentered an unrelated update while holding the semaphore. Patch04 detects an absent target and runs

the original transformation on the same captured solution after releasing protection, preserving errors/exception handling. All 14 fixed error cases then complete. These defects belong to our patches.

RQ2 checks linked propagation, event payloads and snapshots. Variable/lazy work, allocation and VersionStamp generation limit transfer to counts. Custom loaders, changed-text notification reentry and arbitrary compiler behavior remain outside scope. No upstream numerical call cap or latency SLA was found; r is authored. A stronger prepared-rebase comparator reuses existing WithDocumentContentsFrom when solution ID, related-document IDs and all affected project states remain unchanged. It merges prepared contents into the latest solution under protection, otherwise using three-mode fallback. This partial operation lies outside our whole-kernel interface.

7 Implementation and Evaluation

The evaluation asks whether implementations reproduce the resource and information distinctions (RQ1), and which source semantics, resource counts and timing effects the integration supports (RQ2). All studies are authored exploratory tests, with fixed input generation before their recorded executions; none is held-out application data. A later analysis of saved arrivals is explicitly post-outcome. Old failures, timeouts, invalids and corrections remain in the artifact.

An Apple M2/24 GiB runs Python 3.10.9/3.12.14 and OpenJDK 17.0.19; inspected map source matches the cited version but its binary was not rebuilt. Hashes bind inputs, code, commands and traces. Codex assisted design, checkers, source comparison, proofs, analysis and writing; Claude supplied separate proof/direction critiques and manuscript reviews. Measurements are archived executions. These are author-side checks, not mechanical proofs, independent third-party reproduction or human studies.

An *input* fixes a graph, weights and relevant prices/footprints. A *root* is one starting game with specified budget and retry allowance. A recorded native *maximum* is over its fixed finite environment set; the theorems, rather than finite enumeration, justify the unrestricted worst-case statement. Table 2 distinguishes these obligations and does not sum overlapping studies.

7.1 RQ1: Do the Resource and Information Boundaries Hold?

Table 2. Checks relevant to the central resource claims. Each row has its own denominator; some earlier inputs recur in later extensions. Separate code reconstructs the compared quantities, while the proofs establish the general claims.

Obligation	Evidence and comparison
Whole-kernel frontier	Direct equations on 265,629 budget/slack pairs from 5,421 DAGs. Native aware/unaware selectors: 9,314/17,714 records, 1,024/1,280 finite maxima.
Local/shared and fixed toll	Direct games: 35,796 local/shared-cap roots and 42,084 call-toll roots; canonical policies and separate bounds agree.
Single-job repair/toll	Repair: 3,495 inputs/160,080 checks. Per-call toll: 143,226 known-budget values and 33,688 competitive checks. Exact dirty-set games compare with curve/certificate methods.
Shared-retry workflows	8,960 graph/footprint/slack conditions, 61,246 known-budget roots. A quantified one-policy search on 6,208 conditions agrees with the all-budget existence criterion.
Call-charged workflows	37,248 conditions/216,282 budget roots, with a separate one-policy existence search on every condition.

The native unaware selector receives no writer-budget, price or table state; tested budgets reach four. Separate counters reconstruct W, L, Q . A separate suffix-existence recursion agrees on 665,500 states over 363 weight tuples, with 20,594 native paths and 384 finite maxima. Certified post-threshold rules preserve the frontier, but a separate 176,640-root all-tail study gives only 285 strict work improvements and 176,355 ties. The artifact distinguishes this denominator from its earlier 176,608-root adaptive-threshold study. A larger 6,544-input comparison makes greedy certified choice better/worse/tied against all-tail at 1/3/209,404 roots. Constructed counterexamples refute global W -optimality of both rules and simultaneous optimality from every intermediate state. They do not refute suffix existence or the archived threshold-family optimum.

For the new shared-retry study, all ordered pairs of 24 one-/two-component footprint types cover an empty DAG and either chain direction, with $r = 0, 1, 2, 3$. Ordered triples of four representative types cover all eight forward-edge DAGs. The known-budget game chooses a ready job, observes a dirty set, then chooses acceptance or rejection. A separate existence search quantifies one policy over all dirty histories, using accumulated work and minimum consistent writes; it does not read the theorem's saturation predicate as its answer. All 8,960 conditions and 6,208 existence tests agree, including the nonsaturated two-job separation. Four corrupted profiles are rejected. The positive-toll extension checks all 37,248 common-policy conditions, including four further corrupted profiles and seven cases where conditional rejection is required. All initial execution failures and corrections are retained in the artifact. This is bounded corroboration of the proof and its information pattern, not a new application population.

Direct action-before-comparison games also agree on all 35,796 local/shared-cap roots, including 8,928 known-budget local-cap comparisons, and 42,084 fixed-toll roots. Each study detects four changed-input/result controls. These games cover serial fresh preparations; the proofs supply the broader program-class lower bounds.

The new threshold representation matches all 37,248 observed existence answers and 120 additional three-component profile cases. A checker importing no constructor verifies 123 exported tables from full dirty-set observations and rejects four changed certificates; exhaustive executable traces check two common policies. The table uses 803,712 states over the earlier grid versus the boolean search's 616,039: eliminating accumulated-cost enumeration does not ensure fewer states on these small inputs. Partial observations can require more writes than full saturation; only the comparator coordinate is clipped.

The repair compiler retains all 21,028 small/rational comparisons with direct games. Six initial scale timeouts, including two compact-method timeouts, remain; a changed bucket search completes its fixed 168-unit successor. Additive-excess checks cover 4,220 outcomes. These results assess supplied mathematical footprints; they do not calibrate Roslyn merges or establish that arbitrary footprint profiles are inexpensive.

A controlled operational consequence. A Java probe uses four dependent kernels $(1, 2, 3, 5)s$, $r = 0, 1, 2$, and exactly $B = r$ initial failures. Both unaware policies have $Q = 4 + r$, $W = (11 + r)s$; two/three modes protect $11s/0$ iterations. A budget-aware two-mode policy could also avoid protection. During every completing callback, a background compute on another key signals readiness before entry; the foreground holds its bin until that signal. Shared/disjoint-bin layouts have the same hook. Four scales, three allowances, two layouts/modes and three JVMs give 4,320 measurements/1,440 warmups; all complete and match 12 output references. A cross-job checker rejects 12 corruptions. At the largest scale, summed waits fall from 26.8 ms to $31.2\mu s$ in shared bins. Disjoint differences are submicrosecond and sometimes reverse. All 24 cells remain. This forced collateral-blocking probe is not natural-arrival performance or a latency bound.

7.2 RQ2: What Does the Source Integration Preserve?

Source projection and semantic checks. Pinned net8 Roslyn assemblies use a pinned .NET 10 RC runtime. Four projects' evaluated `Compile` items yield 3,837 files, 4,019 document occurrences and 40.8 MB, excluding two package-external items before outcomes. Fixed targets maximize sharing then size (63,593 bytes, two projects) and single-project size (background: 977,535 bytes). This source projection has internal references; updates append comments.

The initial 174-case grid covers $r = 0, 1, 2$, concentrated/distributed writers, same/linked targets, no-ops, distinct equal-content texts and missing/deleted targets. Patch04 plus four unmodified controls pass: 1,645 events, 941 publications, 1,975 snapshots and ten rejected corruptions. A separate checker binds immediate/queued payloads and callbacks to their solutions, checks publication chains and linked transitions, and preserves 4,016 untouched text identities/versions and checked metadata.

Adding 58 rebase cases yields 236 passing cases, 2,190 events, 1,253 publications, 2,653 snapshots and 12 rejected corruptions. Its $B = 5, 6$ cases concentrate initial writes to realize B within fewer calls: placement changed before outcomes; interleavings differ. Error-reentry/same-text controls pass 18/5 cases versus patch04's 14/4. Earlier 204 patched/16 unmodified scale checks, initial 3,960 measurements, same-text deadlock and four missing-document timeouts remain. These projections do not establish full concurrent equivalence. A 96-process study changes types/constants, unrelated source, parse options, references or linked text across eight arms and cold/warm caches. All 2,556 project projections match declared source, symbol, diagnostic and emission expectations; eight corruptions reject. Workspace/direct compilations share Roslyn; expectations cover only these authored programs.

What earlier measurements establish. Earlier forced-interference patch03 runs passed 990 measured/198 warm-up units without later payload validation; three-mode foreground medians improved in four settings and worsened in seven. An independent-arrival successor served 32 serial updates and 64 FIFO background requests at intervals $d \in \{0, 10, 100, 1000\} \mu s$. A six-block balanced campaign passed 1,728 units, and a later rebase comparison passed 3,072 units. Texts were precomputed, the pre-lock hook was null, and no writer bound was supplied to the policy.

These runs show count changes and mixed timing. At shared $r = 2$, two/three modes used 10,802/10,786 calls over 320 measured batches and performed 7,131/3,641 full protected transformations. Three modes also increased total transformations to 14,427. At $d = 100 \mu s$, background p95 medians were 60.00/306.10 μs for three/two modes, lower in seven of eight blocks; foreground medians had mixed block directions. At $d = 0$, three-mode foreground time was 993.54 versus 603.88 μs . No consistent general speedup follows.

Prepared rebase uses stronger source operations. Eligible unrelated-project writes lead to zero full protected transformations, but 4,219/4,528 protected merges remain at $r = 0/2$. A native multi-document batch version passes 236 semantic, 23 reentry and 96 compiler cases. Its 1,728 additional arrival units have batch/sequential median ratios 0.906–1.189 for foreground and 0.511–1.404 for background p95. Protected-span ratios 0.917–1.281 exclude acquisition/release. A document-state guard then changes four reference-removal decisions from recomputation to merge while all 96 compiler outcomes agree; its timing is unmeasured. None of these operations is assumed to have the whole-kernel or additive repair cost model.

Comparing shared and per-update contracts. The local-cap baseline resets the existing retry field before each update; the checked patch04 libraries are unchanged. We fix eight arms: unmodified baseline, instrumented original, shared two/three at $r = 2$, and per-job two/three at $r = 1, 2$. Eight cyclic Latin blocks place every arm at every process position. Each process has two warm-up




Fig. 2. New shared-cap timing comparison, $r = 2$. Points are geometric means of eight paired block ratios of process medians; bars are 95% percentile intervals from 10,000 fixed block-bootstrap resamples. Ratios above 1 favor two modes. Writer response is intended arrival to return; service is invocation to return. Each batch contributes its median across 64 writer requests. All four intervals and three metrics are shown.

and ten measured batches at each of the four intervals, shuffled before execution. All 3,072 units (2,560 measured/512 warm-up) complete and pass the source checks; all 12 corruption controls are detected. This is a new authored workload comparison, not a new deployment population.

The additional per-update deltas sum to each batch's counts. For an update's actual publication count B_i , local modes satisfy $Q_i = 1 + f_i$, $f_i \leq \min(B_i, r)$, and inside count at most $1[B_i \geq r]$ (two) or $1[B_i \geq r + 1]$ (three). With B_{fg} counting publications between foreground start and end, summing gives the uniform count bounds of Proposition 3. All individual updates and batches meet their own contract. Table 3 shows every arrival interval; each entry contains separate medians, which need not coexist on a single execution.

Table 3. New Roslyn cap comparison, 80 measured batches per cell. Entries are median $Q/B_{fg}/N_{inside}$: counted calls, foreground-period writes and full protected transformations. The universal call cap differs across rows.

Policy	Q cap	$d = 0$	10	100	1000 μs
Shared two, $r = 2$	34	34/3/32	34/3/32	34/2/31	33/1/0
Shared three, $r = 2$	34	34/34/30	34/27/22	34/8/6	33/1/0
Local two, $r = 1$	64	62/31.5/30	54/24/22	40/9/8	33/1/1
Local three, $r = 1$	64	58/64/25	53/50.5/19	37/7/1	33/1/0
Local two, $r = 2$	96	79.5/64/23.5	77/57/22	39/7.5/1	33/1/0
Local three, $r = 2$	96	66/64/17	64.5/64/16	38/7/0	33/1/0

The local baselines can protect fewer updates, while allowing more calls. For example, at 100 μs the local three-mode $r = 2$ medians are 38 calls and zero protected transformations; shared $r = 2$ has 34 and 6. Their caps are 96 and 34, respectively. This illustrates a resource-contract choice, not dominance under equal constraints. The actual B_{fg} depends on the policy's completion time. A separate post-outcome reconstruction of 7,392 earlier batches also verifies that coordinate and all recorded count bounds; it does not turn the comparisons into a common controlled write budget.

For shared three/two, foreground ratios at 0/10 μs are 1.72 [1.40,2.11] and 1.58 [1.17,2.16]; the two slower-arrival intervals include 1. All four writer-response intervals include 1 (Figure 2). Added per-update counter snapshots occur inside foreground batch duration. Original/baseline point ratios are 0.83–1.08 and all four intervals include 1, so this run does not identify instrumentation overhead; earlier harness estimates remain separate. Block resampling does not remove residual

host/order effects, and no multiplicity correction is applied. Native execution took 237.84 seconds after the prior public replay completed, and the artifact includes setup/build and analysis costs. These results connect cap scope to checked counts while limiting any timing interpretation.

Earlier auxiliary objectives. The artifact retains separate supplied-price compilation, calibration and Deephaven studies, with their complete definitions and all outcomes. They do not evaluate the common-policy repair interface. Three-mode scalar gains occurred in only 2/96 cells; a time fit met its fixed residual threshold in only 2/8 settings. General all-budget compilation sometimes cost more than requested-value computation, and a Deephaven phase-aware model improved a previous abstraction's value. These adverse results motivate keeping resource guarantees separate from timing or an exact native reduction.

8 Related Work and Research Lineage

Fallback and online contention management. Kung and Robinson [3], Section 3.3/p.220, already retain the critical-section semaphore during reexecution after repeated validation failure. Lock-elision retry caps likewise precede our work [38, 39]. Cached completion here requires an unrolled-back preparation and atomic identity-check/recompute; HTM capacity or explicit aborts need not consume a write [40]. These mechanisms motivate the interfaces but do not establish their joint Q, L guarantee.

Attiya et al. [36], Sections 2–3/5, analyze one Greedy contention manager that does not receive the per-job non-conflict failure bound k . It is $O(ks)$ -competitive for makespan with s resources, with a corresponding deterministic lower bound under time-dependent resource requests. Thus a single policy for an unknown disturbance level and an all-policy lower bound are established. Our allowance counts global external writes, the cap constrains every finite-write execution, and the output is an exact protected-kernel-work curve for persistent jobs. Alistarh et al. [37], Sections 4–5, optimize whether and when to abort a conflicting transaction with known abort cost and unknown remaining duration. Their original throughput policy has no progress guarantee; Section 7 adds a probabilistic one. That objective and returned guarantee differ from the universal call cap here.

Guerraoui, Herlihy and Pochon [1], Sections 3–4, establish a priority-based abort/wait manager with bounded-delay progress and a makespan guarantee against an offline scheduler. Diegues and Romano [2], Sections 4–6 and its errata, adapt attempt budgets and capacity-abort handling using per-atomic-block measurements in GCC. Thus theoretically justified contention control and workload-oblivious adaptive fallback are prior contributions. Their transaction progress/cycle objectives differ from the persistent-preparation operations and hard local/shared count contracts here; their measurements do not calibrate our call toll.

Synchronization synthesis and robust control. Černý et al. [20] choose retries, guards and caches in partial concurrent programs with a scheduler and weighted limit-average safety objective. Their Theorem 6 uses quantitative probabilistic bisimulation; the erratum corrects complexity to PSPACE. Later work optimizes protected-statement counts and statistically profiles acquisition/contention [21]; preemption-safety synthesis also addresses deadlock-free lock placement [22]. Our archived scalar potential concerns initial worst total cost, not source equivalence or intermediate-state quotients. Finite n, B, r instances can be encoded as general games; we do not claim such tools cannot express them.

Mannor et al. [11] supply global-deviation-budget DP; Postek et al. [31] provide finite-candidate DP and causal budgeted routing. Information-state minimax control [41] is a general framework for the repair stopping problem after proving its least-consistent-budget reduction. Our extra repair argument uses independently realizable intervals and monotonicity to reduce a feasibility tree to

one minimum-write bad path. Generic backward induction and sufficient information are not new claims.

Inherited auxiliary results. The archived priced analysis uses reflected recurrences, robust envelopes, concave allocation and composition ordering [7, 15, 16, 29, 32]. Fault-tolerant scheduling, resilient synthesis, decision-dependent optimization and ski rental have different timing or information contracts [12–14, 17–19, 23, 24, 28, 33]; a different recurrence alone does not preclude a reduction.

The zero-fee snapshot [34] is an *unpublished, unreviewed author artifact*. Its floor/cap lemma, two-fee ordering and restricted hardness concern supplied-price compilation. These auxiliary results and all adverse outcomes remain archived; the present theorems concern resource frontiers and partial-repair information boundaries.

9 Limitations and Threats to Validity

The exact resource frontier concerns mandatory whole kernels, persistent outputs and the listed completing calls. It does not include replenished write budgets, structural change, exceptions, external completion or retained guards. In retained charged-guard studies, five of 384 inputs improve over serial calls (including 242 versus 243); an intermediate state separates immediate completion/release from retaining the guard (43 versus 45). General retained-guard equality remains unproved. These counterexamples and the repair theorems delimit the result rather than refute a theorem under its stated interface.

Roslyn experiments cover selected source paths and authored changes. Same-text deadlocks and missing-document timeouts show why replayability does not imply safe protected execution. Event, snapshot and compiler projections do not establish arbitrary concurrent equivalence; direct/workspace compilations share the compiler engine. Rebase retains merge cost and is outside whole-kernel optimality. Our component profiles and call tolls have not been calibrated to that merge implementation. Footprint profiling may be exponential, and native validation can reveal different information.

All measurements use one author workstation, authored arrivals and recorded process variability. Instrumentation costs and cache/JIT effects prevent inference of uninstrumented or production timing; no human benefit is measured. Scalar gains are sparse, the attempted time calibration fails in most settings, and requested-value methods can cost less than a compiler. Finite equation tests and AI-assisted proof inspection are not mechanical proofs or independent third-party reproductions.

10 Conclusion

A finite retry allowance does not by itself determine protected work. Under the whole-kernel completion interface, cached reuse/recompute permits one policy to meet a universal call cap and the exact protection frontier for every hidden write bound. Partial repair and per-call costs have different, sharp information boundaries. These characterizations specify when an interface supports the desired resource contract. Source checks and arrival experiments expose both realizable count tradeoffs and the limits of interpreting them as timing improvements.

Data Availability

The public supplement [35] contains proof details, fixed inputs, sources, raw outcomes, corrections and replay commands; the zero-fee evidence [34] is unchanged. The central resource and repair theorems are proved in this paper; archived auxiliary priced results are not needed for them. Original/public hashes record path projections. Licensed modified sources retain notices; upstream archives, runtimes, caches, articles and private correspondence are excluded.

References

- [1] Rachid Guerraoui, Maurice Herlihy, and Bastian Pochon. 2005. Toward a Theory of Transactional Contention Managers. In *PODC*. Author version. <https://lpdwww.epfl.ch/upload/documents/publications/neg--509904679f181-pochon.pdf>.
- [2] Nuno Diegues and Paolo Romano. 2014. Self-Tuning Intel Transactional Synchronization Extensions. In *ICAC*, 209–219. <https://www.usenix.org/conference/icac14/technical-sessions/presentation/diegues>. Errata.
- [3] H. T. Kung and John T. Robinson. 1981. On optimistic methods for concurrency control. *ACM Transactions on Database Systems* 6, 2, 213–226. <https://doi.org/10.1145/319566.319567>.
- [4] .NET Foundation. 2025. *Roslyn source*, commit 6c4a46a31302167b425d5e0a31ea83c9a9aa1d09. `Workspace.cs`, `Solution.cs`, `SolutionCompilationState.cs` and `SolutionState.cs`. Accessed September 10, 2026. <https://github.com/dotnet/roslyn/tree/6c4a46a31302167b425d5e0a31ea83c9a9aa1d09>.
- [5] .NET Foundation contributors. 2020. Add `SetCurrentSolution` with solution transformation, pull request 42228. Source review on replayability and event ordering. Accessed September 10, 2026. <https://github.com/dotnet/roslyn/pull/42228>.
- [6] OpenJDK. n.d. *JMH: Java Microbenchmark Harness*, version 1.37. Accessed September 9, 2026. <https://github.com/openjdk/jmh/tree/1.37>.
- [7] Yasushi Kawase, Kazuhisa Makino, and Kento Seimi. 2016. Optimal Composition Ordering Problems for Piecewise Linear Functions. In *ISAAC, LIPIcs* 64, 42:1–42:13. <https://doi.org/10.4230/LIPIcs.ISAAC.2016.42>. Extended version: *Algorithmica* (2018), <https://doi.org/10.1007/s00453-017-0397-y>.
- [8] Deephaven Data Labs. 2026. *Deephaven Core source*, commit 6367313a319b79437d4d2116c77836ad56d5138d. `Files ConstructSnapshot.java`, `HierarchicalTableImpl.java`, `EventDrivenUpdateGraph.java`, and `JsTreeTable.java`. Accessed September 9, 2026. <https://github.com/deephaven/deephaven-core/tree/6367313a319b79437d4d2116c77836ad56d5138d>.
- [9] Deephaven Data Labs. 2026. *tree: Real-time tree table*. Python documentation, version 0.40.7. Accessed September 9, 2026. <https://deephaven.io/core/0.40.7/docs/reference/table-operations/create/tree/>.
- [10] Deephaven Data Labs. 2026. *Execution context*. Official documentation, custom update graphs. Accessed September 9, 2026. <https://deephaven.io/core/docs/conceptual/execution-context/>.
- [11] Shie Mannor, Ofir Mebel, and Huan Xu. 2012. Lightning Does Not Strike Twice: Robust MDPs with Coupled Uncertainty. In *ICML*. <https://icml.cc/2012/papers/215.pdf>.
- [12] Daniel Neider, Alexander Weinert, and Martin Zimmermann. 2020. Synthesizing optimally resilient controllers. *Acta Informatica* 57, 195–221. <https://doi.org/10.1007/s00236-019-00345-7>.
- [13] Alexander Weinert. 2019. *Optimality and Resilience in Parity Games*. PhD thesis, Saarland University. Section 6.2. https://finkbeiner.groups.cispa.de/_assets/paper.C89PMq4G.pdf.
- [14] Hyung-Chan An and Hoesook Yang. 2021. Fully Adaptive Stochastic Handling of Soft-Errors in Real-Time Systems. *IEEE Access* 9, 155058–155071. <https://doi.org/10.1109/ACCESS.2021.3128397>.
- [15] Awi Federgruen and Henri Groenevelt. 1986. The Greedy Procedure for Resource Allocation Problems: Necessary and Sufficient Conditions for Optimality. *Operations Research*. https://business.columbia.edu/sites/default/files-efs/pubfiles/4071/federgruen_greedy_procedure.pdf.
- [16] Dimitris Bertsimas and Melvyn Sim. 2004. The Price of Robustness. *Operations Research* 52, 1, 35–53. <https://www.mit.edu/~dbertsim/papers/Robust%20Optimization/The%20price%20of%20Robustness.pdf>.
- [17] Marek Chrobak, Mathilde Hurand, and Jiří Sgall. 2009. Algorithms for Testing Fault-Tolerance of Sequenced Jobs. *Journal of Scheduling*. Author version. <https://iuuk.mff.cuni.cz/~sgall/ps/faults.pdf>.
- [18] Phebe Vayanos, Angelos Georgioui, and Han Yu. 2026. Robust Optimization with Decision-Dependent Information Discovery. *Management Science* 72, 2, 1509–1528. <https://doi.org/10.1287/mnsc.2021.00160>.
- [19] Eva Ley, Rebecca Marx, Maximilian Merkert, Tim Niemann, and Sebastian Stiller. 2026. Decision-dependent robust optimization with elementwise controllable uncertainty sets. *OR Spectrum*. <https://doi.org/10.1007/s00291-026-00856-z>.
- [20] Pavol Černý, Krishnendu Chatterjee, Thomas A. Henzinger, Arjun Radhakrishna, and Rohit Singh. 2011. Quantitative Synthesis for Concurrent Programs. In *CAV*, 243–259. https://doi.org/10.1007/978-3-642-22110-1_20. Author publication page and complexity erratum: <https://www.microsoft.com/en-us/research/publication/quantitative-synthesis-for-concurrent-programs/>.
- [21] Pavol Černý, Edmund M. Clarke, Thomas A. Henzinger, Arjun Radhakrishna, Leonid Ryzhyk, Roopsha Samanta, and Thorsten Tarrach. 2015. Optimizing Solution Quality in Synchronization Synthesis. *arXiv preprint arXiv:1511.07163v1*. <https://arxiv.org/abs/1511.07163>.
- [22] Pavol Černý, Edmund M. Clarke, Thomas A. Henzinger, Arjun Radhakrishna, Leonid Ryzhyk, Roopsha Samanta, and Thorsten Tarrach. 2017. From non-preemptive to preemptive scheduling using synchronization synthesis. *Formal Methods in System Design* 50, 97–139. <https://doi.org/10.1007/s10703-016-0256-5>.
- [23] Viacheslav Izosimov, Paul Pop, Petru Eles, and Zebo Peng. 2012. Scheduling and Optimization of Fault-Tolerant Embedded Systems with Transparency/Performance Trade-Offs. *ACM Transactions on Embedded Computing Systems* 11, 3, Article 61, 35 pages. <https://doi.org/10.1145/2345770.2345773>.

- [24] Bingham Wu, Wei Bao, Dong Yuan, and Bing Zhou. 2022. Competitive Analysis for Multi-Commodity Ski-Rental Problem. In *IJCAI*, 4672–4678. <https://doi.org/10.24963/ijcai.2022/648>.
- [25] Oracle. 2021. ConcurrentHashMap, Java SE 17 API Specification. Accessed September 7, 2026. <https://docs.oracle.com/en/java/javase/17/docs/api/java.base/java/util/concurrent/ConcurrentHashMap.html>.
- [26] Oracle. 2021. ConcurrentMap, Java SE 17 API Specification. Accessed September 7, 2026. <https://docs.oracle.com/en/java/javase/17/docs/api/java.base/java/util/concurrent/ConcurrentMap.html>.
- [27] OpenJDK. 2026. ConcurrentHashMap.java, jdk17u source, commit 982d53c227943a511e8d61694feced764b71a6c1. Accessed September 7, 2026. <https://github.com/openjdk/jdk17u/blob/982d53c227943a511e8d61694feced764b71a6c1/src/java.base/share/classes/java/util/concurrent/ConcurrentHashMap.java>.
- [28] Ziwei Li, Bo Sun, Zhiqiu Zhang, Mohammad Hajiesmaili, Bingham Wu, Lin F. Yang, and Yang Gao. 2025. Combinatorial Ski Rental Problem: Robust and Learning-Augmented Algorithms. In *NeurIPS*. https://papers.nips.cc/paper_files/paper/2025/hash/ae3e173398d6e43fea63cbcc16fbaa98-Abstract-Conference.html.
- [29] Krzysztof Burdzy, Weining Kang, and Kavita Ramanan. 2009. The Skorokhod problem in a time-dependent interval. *Stochastic Processes and their Applications* 119, 2, 428–452. <https://doi.org/10.1016/j.spa.2008.03.001>.
- [30] Laurent Perron, Frédéric Didier, and Steven Gay. 2023. The CP-SAT-LP Solver (Invited Talk). In *CP, LIPIcs* 280, 3:1–3:2. <https://doi.org/10.4230/LIPIcs.CP.2023.3>.
- [31] Krzysztof Postek, Ward Romeijnnders, and Wolfram Wiesemann. 2025. Technical Note: Multi-Stage Robust Mixed-Integer Programming. *Operations Research* 73, 6, 3345–3358. <https://doi.org/10.1287/opre.2023.0520>. Online appendix; author code.
- [32] Dimitris Bertsimas and Melvyn Sim. 2003. Robust discrete optimization and network flows. *Mathematical Programming* 98, 49–71. <https://doi.org/10.1007/s10107-003-0396-4>.
- [33] Matthew Bold and Marc Goerigk. 2021. A compact reformulation of the two-stage robust resource-constrained project scheduling problem. *Computers & Operations Research*. <https://doi.org/10.1016/j.cor.2021.105232>. Accepted manuscript.
- [34] Anonymous Authors. 2026. *Compiling Optimal Retry Policies for Dependent Transformations: Research Artifact*. Version 5 with publication supplement, commit e0e0110. <https://github.com/research-artifact-archive/research-artifact/tree/e0e011023c0591bbbf422390ad8d29785ad0023f>.
- [35] Anonymous Authors. 2026. *Retry Synthesis with Bounded Calls and Protected Work: Research Artifact*. Version 18 local/shared caps and per-update comparison, commit 7e93190. <https://github.com/research-artifact-archive/research-artifact/tree/7e93190645f2f16f21539ad8bcb32404722ca294>.
- [36] Hagit Attiya, Leah Epstein, Hadas Shachnai, and Tami Tamir. 2006. Transactional Contention Management as a Non-Clairvoyant Scheduling Problem. In *PODC*. <https://cs.idc.ac.il/~tami/Papers/podc06.pdf>.
- [37] Dan Alistarh, Syed Kamran Haider, Raphael Kübler, and Giorgi Nadiradze. 2018. The Transactional Conflict Problem. *arXiv preprint arXiv:1804.00947v1*. <https://arxiv.org/abs/1804.00947>.
- [38] Free Software Foundation. 2025. GNU C Library 2.42: x86 lock elision, revision d2097651. *elision-lock.c* and *elision-conf.c*. <https://github.com/bminor/glibc/tree/d2097651cc57834dbfcaa102ddfacae0d86cfb66/sysdeps/unix/sysv/linux/x86>. Official tunables: https://www.gnu.org/software/libc/manual/html_node/Elision-Tunables.html.
- [39] OpenJDK. 2025. HotSpot 21.0.8+9 RTM monitor code, revision 54f20959. *c2_MacroAssembler_x86.cpp* and *globals_x86.hpp*. <https://github.com/openjdk/jdk21u/tree/54f2095960f01d957f2335cafa0defb956e13a2c/src/hotspot/cpu/x86>.
- [40] GNU Compiler Collection. n.d. x86 Transactional Memory Intrinsics. Accessed September 10, 2026. <https://gcc.gnu.org/onlinedocs/gcc/x86-transactional-memory-intrinsics.html>.
- [41] Aditya Dave, Nishanth Venkatesh, and Andreas A. Malikopoulos. 2024. Approximate Information States for Worst-Case Control and Learning in Uncertain Systems. *arXiv preprint arXiv:2301.05089v2*. <https://arxiv.org/abs/2301.05089>.